Untitled: A DirectX Game

Sam Drysdale

May 16, 2023

Contents

1 Summary

Untitled is [...].

2   User Controls

3   Features

The following is a technical discussion of the key features of *Untitled*, with a focus on advanced procedural generation. Mathematical models and code snippets are kept to a minimum, edited for clarity rather than accuracy to the original application.

3.1   Noise

3.2   Procedural Terrain

Given its grounding in nuclear semiotics, this game's terrain is of course intended to evoke a sense of "shunned land" (**?**). In generating the jutting thorns and blocks so essential to the post-nuclear setting, though, there comes a problem - concavity. Height maps are perfectly adequate for creating rolling hills and valleys, but the fact that each (x, z)-coordinate can correspond to only one y-value prevents any features that ominously overhang *Untitled*'s barren earth (see Figure **??**).


Figure 1: Two sketches of *Untitled*'s thorny terrain; though the former can be generated by a height map, the latter is concave.

Consider the two-dimensional analogue of this problem: how does one construct the bounding curve of a (potentially concave, or even disconnected) area? Taking an $(n + 1) \times (n + 1)$ grid, with $n \in \mathbb{N}$,[1] let the scalar field $f : [0, 1]^2 \to \mathbb{R}$ assign each gridpoint (i, j) a corresponding value $f_{i,j} = f(i/n, j/n)$. *Marching squares* can then be used to approximate the *isoclines* - or informally, the 'contour lines' - $f(x, y) = c$ constant. [By partitioning each of the $n \times n$ square cells according to which of its four gridpoints have $f_{i,j} > c$ (see Figure **??**), an increasingly accurate curve emerges as $n \to \infty$.]

Figure 2: Partitions of a marching square. With each of the 4 gridpoints existing in one of 2 states ('+' for $f_{i,j} > c$, '−' for $f_{i,j} \leq c$), exactly $2^4 = 16$ partitions are possible.

---

[1]This could equally be an $(n + 1) \times (m + 1)$ grid; keeping the dimensions the same is just 'neater'.

[Note that, with... Figure **??**...]

$$f(x, y) = \min\left((x - 0.5)^2 + (y - 0.5)^2, 4\left((x - 0.75)^2 + (y - 0.75)^2\right)\right),$$

a scalar field describing two overlapping circles.

Figure 3: Bounding curves of $f(x, y) = 0.16$, at increasing levels of granularity.

[Linear interpolation] Though . Again, this is just an approximation, but the improvement should be immediately apparent (see Figure **??**).


Figure 4: Bounding curves of $f(x, y) = 0.16$, before and after linear intepolation.

In much the same way, *Untitled* uses *marching cubes* to generate *isosurfaces* bounding 3D volumes. Outside of the expected changes - scalar fields now operate on domain $[0, 1]^3$, vertices now have $2^8 = 256$ possible configurations along [the edges of] each marching cube - this is actually a rather straightforward generalisation.

[These are just the broad strokes; nuance to, say, weighing vertex normals correctly]. Definitive... Paul Bourke's *Polygonising a Scalar Field* (**?**) remains the definitive source on the matter.

3.2.1 Case Study: Hexes

Having . 3D fields are inherently , yet these are essential to carving out *Untitled*'s landscape.

Even a blank hex tile comes with som nuance. [Use of noise...]

[Bounding hexagonal prism... As much as flat faces are antithetical to marching cubes (if they run parallel to the grid, then... can't interpolate? notice that interpolation doesn't account for the degenerate case where two adjacent gridpoints are equal)...]

3.2.2 Case Study: Landmarks

3.3 Procedural Screen Textures

The post-processing in *Untitled* is, in one sense, rather simple. The 'stress vignette,' for instance, calls only two renders-to-texture on every frame: the board itself, and an alpha map of blood vessels that sprout from the edges of the screen. As striking as the final effect is, `vignette_ps.hlsl` is surprisingly straightforward in blending the textures into a single, pulsing eye strain overlay; far more deserving of further discussion is how the blood vessels themselves are generated.

In formal languages, a grammar is a tuple $G = (N, \Sigma, P, \omega_0)$. This contains two disjoint sets of symbols: nonterminals $A, B, \cdots \in N$, and terminals $a, b, \cdots \in \Sigma$. The production rules in P map nonterminals to strings $\alpha, \beta, \cdots \in (N \cup \Sigma)^*$; applied recursively to the axiom $\omega_0 \in (N \cup \Sigma)^*$, these rules can produce increasingly complex *sentences* of terminals and/or nonterminals.[2]

The Chomsky hierarchy (**?**) classifies grammars by their production rules:

> *Type-3. Regular grammars* map $A \mapsto a$ or $A \mapsto aB$.
>
> *Type-2. Context-free grammars* map $A \mapsto \alpha$.
>
> *Type-1. Context-sensitive grammars* $\alpha A \beta \mapsto \alpha \gamma \beta$.
>
> *Type-0. Unrestricted grammars* map $\alpha \mapsto \beta$, where α is non-empty.

Note that all Type-3 grammars are also Type-2, all Type-2 grammars also Type-1, and so on.

Suppose, for example, that $N = \{F, G\}$, $\Sigma = \{+, -\}$, $P = \{F \mapsto F + G, G \mapsto F - G\}$, $\omega_0 = F$. Letting ω_n denote the sentences generated by applying the production rules n times, it follows that

$$
\begin{aligned}
\omega_1 &= F + G, \\
\omega_2 &= F + G + F - G, \\
\omega_3 &= F + G + F - G + F + G - F - G, \\
\omega_4 &= F + G + F - G + F + G - F - G + F + G + F - G - F + G - F - G, \cdots .
\end{aligned}
$$

While these defintions are rather abstract, **?** provides a remarkable application. Treating each symbol as an instruction like 'go forward' or 'turn right', *L-systems* visualise sentences via 'turtle graphics'; when those sentences have been generated recursively by a grammar, the line drawings inherit that same self-similar structure. Consider the above, where reading non-terminals F, G as 'draw a line while moving one unit forwards,' and terminals $\pm$ as 'turn $\pm \pi/2$ on the spot,' produces the fractal dragon curves in Figure **??**.

While L-systems are most common in the modelling of 3D plants and other branching structures (**?**), *Untitled* only uses them to generate 2D alpha maps. Moreover, it restricts its attention to L-systems paired with context-free grammars.

---

[2]In mathematical literature, $\omega_0 \in N$ (**?**), but *Untitled* takes an informal approach.

Figure 5: Dragon curves, generated by strings $\omega_2, \omega_4, \cdots, \omega_{12}$.

3.3.1 Case Study: Runes

Parametric L-systems (**?**) exist as a generalisation of the above. [theory].

The modules in *Untitled*, then, track three . More than anything, this offers a certain clarity of code - at least from a

[Example: various geometric runes!].

3.3.2 Case Study: Blood Vessels

?, meanwhile, uses parametric L-systems to visualise the bifurcation of blood vessels. Suppose a branch with length l, width w bifurcates into two branches M and m, such that $l_M \geq l_m$. Defining the *asymmetry ratio* $\alpha = l_m/l_M$, it follows that

$$l_M = \frac{l}{(1+\alpha^3)^{1/3}}, \quad l_m = \frac{\alpha \cdot l}{(1+\alpha^3)^{1/3}}, \quad w_M = \frac{w}{(1+\alpha^3)^{1/3}}, \quad w_m = \frac{\alpha \cdot w}{(1+\alpha^3)^{1/3}}.$$

Furthermore, the branches diverge from their parent at angles

$$\theta_M = \arccos\left(\frac{(1+\alpha^3)^{4/3}+1-\alpha^4}{2(1+\alpha^3)^{2/3}}\right), \quad \theta_m = \arccos\left(\frac{(1+\alpha^3)^{4/3}+\alpha^4-1}{2\alpha^2(1+\alpha^3)^{2/3}}\right).$$

Untitled's framework is therefore capable of reproducing **?**'s results (see Figure **??**), using an L-system with the single production rule:

$$\mathbf{C}(l,w,\theta) \mapsto \mathbf{X}(l,w,\theta)[\mathbf{C}(l_M,w_M,\theta+\theta_M)]\mathbf{C}(l_m,w_m,\theta-\theta_m).$$

Figure 6: **?**'s model of arterial branching, with asymmetry ratios $\alpha = 1.0, 0.8, \cdots, 0.2$.

? expand on this by introducing a stochastic component - that is to say, they allow [a more random

structure]. *Untitled* incorporates such randomness into its own rules for blood vessels:

$$\mathbf{C}(l, w, \theta) \underset{0.4}{\longmapsto} \mathbf{X}(l, w, \theta)[\mathbf{L}(l_M, w_M, \theta + \theta_M)]\mathbf{R}(l_m, w_m, \theta - \theta_m)$$
$$\mathbf{C}(l, w, \theta) \underset{0.4}{\longmapsto} \mathbf{X}(l, w, \theta)[\mathbf{L}(l_m, w_m, \theta + \theta_m)]\mathbf{R}(l_M, w_M, \theta - \theta_M)$$
$$\mathbf{C}(l, w, \theta) \underset{0.2}{\longmapsto} \mathbf{X}(l, w, \theta)\mathbf{C}(l, w, \theta)$$

$$\mathbf{L}(l, w, \theta) \underset{1.0}{\longmapsto} \mathbf{X}(l, w, \theta)\mathbf{C}(l_M, w_M, \theta - \theta_M)$$

$$\mathbf{R}(l, w, \theta) \underset{1.0}{\longmapsto} \mathbf{X}(l, w, \theta)\mathbf{C}(l_M, w_M, \theta + \theta_M)$$

These describe a capillary with a 40% chance of bifurcating with branch M tacking clockwise, a 40% chance of bifurcating with M tacking anticlockwise, and a 20% chance of extending forwards without any branching. The determinstic production rules on $\mathbf{L}$, $\mathbf{R}$ provide course correction, guaranteeing the [...]; further informal tweaks can be found in the `LBloodVessel` class, all intended to get the final look of the L-systems 'right' (see Figure [reference]).

[Figure of blood vessels in isolation]

[Discussion of animation (and the shortcomings thereof)...]

[Figure of final render]

3.4   Procedural Narrative

Untitled was originally conceived as a showcase of procedural text generation, an application of [authored X] towards interactive fiction.

3.4.1   Grammars

Given their origin in linguistics, it is perhaps unsurprising that grammars (see Section **??**) are of use in the field of procedural narrative - *Tracery* (**?**) being the prime example. From a mathematical perspective, this is a stochastic, context-free grammar, with uniformly-distributed production rules for each non-terminal; it iterates a given axiom until all non-terminals (demarcated by {`SQUARE BRACKET`} delimiters) have been replaced. To bemoan the generator as better suited to Twitter bots than immersive storytelling is to misunderstand its design. **?** present a lightweight tool that is left *deliberately* narrow, intended for ease-of-use amongst even the most fledgling authors.

More substantial works, then, adapt *Tracery*'s grammar-based approach to their own specifications. In *The Annals of the Parrigues*, **?** sets out the need for a consistent generator, one that "defines any facts about the world that aren't already defined at the moment of generation" (**?**, p. 83). *Voyageur* (**?**) attaches weightings to the distributions of production rules (**?**). Produceral storytelling is an art form; as much as *Untitled* wears these influences on its sleeve, to try and consolidate these into an 'optimal', one-size-fits-all generator would be missing the point. This project therefore uses an implementation of *Tracery* with a couple of custom modifications tailored to its own needs.

Recency   [Clear, one sentence goal: avoid unnecessary repetition with recency].[3] *Untitled* defines this by rank: given a non-terminal with N possible production rules, the most recently used rule is will have a recency of $N - 1$, the second most recent $N - 2$, and so on down to the rule that has been used longest ago (or indeed, never been used), with recency 0. Though it [does not consider the

---

[3] *Improv*, the generator for *Voyageur*, calls this same quality 'dryness'.

actual time interval between these uses, nor the *second* most], this metric is already enough to achieve the desired effect (**?**).

[Second paragraph: how does this affect weighting?] $p^{n/(N-1)}$, such that (all other weightings being equal), the most recently used rule (with $n = N - 1$) will be p times as likely as the least recent $(n = 0)$.[4]

Note that [joint recency!].

Characterisation   [Consistency discussion...]

This might feel too limited, from a dramatic perspective. For all this effort to [maintain consistency], personality is not immutable; so much of good storytelling relies on character *development*. Indeed, how does the player feel like they have an impact, if [the grammar doesn't have permission to make changes]? *Untitled* surely requires some sort of override...

3.4.2   Content Selection Architectures

Storylets (**?**) are [definition].

[Though conceptually no different, ... , our 'narrative stack'...]

4   Code Organisation

[...]

Section **??** has been structured [to reflect the two fundamental stages of graphics programming: first implementing frameworks (marching cubes, L-systems), then building tangible objects out of them (terrain hexes, blood vessels)]. *Untitled*'s code is structured accordingly. [Using inheritance...]. [Crucially, little is left exposed in the main *Game.cpp* file...].

4.1   Rendering

4.2   GUI

[Include HDRR/bloom here...]

5   Evaluation

5.1   Features

[Start with marching cubes: successful, but more to do...]

[L-systems: far more robust...]

To the extent that *Untitled* has any single 'special feature', though, it would surely be its approach to procedural storytelling. [...]

5.2   Code Organisation

[Position weakness of the storytelling as an *organisational* matter...]

[Discuss the broader merits of Section 4, ultimately landing on the value of *internal consistency*.]

---

[4]The same string `alpha` might replace two different non-terminals `#A#`, `#B#`. Recognising these as distinct production rules, the recency with which one was applied to `#A#` has no bearing on when the other is applied to `#B#`, or vice versa.

6   Conclusions

From a more personal perspective, I can't help but feel that *Untitled* shows promise.

[Coheres in a way that my CMP502 project absolutely didn't...]. [Furthermore, a better sense of performance] - I might not be happy with the loading times *per se*, but all things considered these advanced graphics techniques will still run on my laptop.

Of course, there is a distinction to be drawn between *Untitled*'s functionality as a graphics showcase and as a game. As much as there's a certain level of interactivity on display,[5] [Perils of interactive narrative].

[What/how would I go about cannibalising this? Screen shader first/hexes... narrative much more an early experiment in structuring content selection architectures/context-sensitive grammars...]

References

Bourke, P. (1994), 'Polygonising a Scalar Field', Available at: `http://paulbourke.net/geometry/polygonise/`. (Accessed: 9 February 2023).

Chomsky, N. (1956), 'Three Models for the Description of Language', *IRE Transactions on Information Theory* **2**(3), 113–124.

Hopcroft, J., Motwani, R. and Ullman, J. D. (2000), *Introduction to Automata Theory, Languages, and Computation*, 2nd edn, Boston, MA, USA: Addison-Wesley.

Lindenmayer, A. (1968), 'Mathematical Models for Cellular Interactions in Development II. Simple and Branching Filaments With Two-Sided Inputs', *Journal of Theoretical Biology* **18**(3), 300–315.

Liu, X., Liu, H., Hao, A. and Zhao, Q. (2010), Simulation of Blood Vessels for Surgery Simulators, *in* '2010 International Conference on Machine Vision and Human-machine Interface', pp. 377–380.

Zamir, M. (2001), 'Arterial Branching Within the Confines of Fractal L-System Formalism', *The Journal of General Physiology* **118**, 267–276.

---

[5]And likewise, non-interactive features like the fixed, orthographic camera carry some level of intentionality...