

Concrete Earth: A DirectX Game

Sam Drysdale

May 16, 2023

Contents

1	Summary	1
2	User Controls	2
3	Features	2
3.1	Procedural Terrain	2
3.1.1	Case Study: Hexes	5
3.1.2	Case Study: Landmarks	6
3.2	Procedural Screen Shaders	6
3.2.1	Case Study: Blood Vessels	8
3.3	Procedural Narrative	9
3.3.1	Grammars	10
3.3.2	Storylets	11
4	Code Organisation	12
5	Evaluation	13
5.1	Features	13
5.2	Code Organisation	14
6	Conclusions	14
	References	15

1 Summary

Concrete Earth is a game about nuclear semiotics. Ostensibly, the story is set in a medieval fiefdom, cursed by the occult ciphers long-etched into its landscape. The twist is this is set millenia *after* a nuclear apocalypse - the ‘curse’ is radiation poisoning, the ciphers the dead language of a previous civilisation, warning of nearby waste repositories. Though it is still in development, this prototype demonstrates how the DirectX 11 functionality might bring such a concept to life. Indeed, for a game so concerned with language, it is appropriate that many key features are grounded in linguistics; while the terrain itself is generated by marching cubes, *Concrete Earth* uses grammars to create both a procedural screen shader and a procedural narrative.

A video demonstration can be found within the “docs” directory, and [online](#).

2 User Controls

The user plays as a carriage driver, journeying forwards over an endless map. Periodically, they will be stopped by a ‘multiple-choice’ narrative event, where an option must be selected before continuing. Such is the core game loop of *Concrete Earth*.

The application uses the following key mappings:¹

W	Move north
A	Move northwest
D	Move northeast
1	Select option #1
2	Select option #2
Esc	Quit

Clicking **LMB** will also select an option. The application runs at a fixed 1920×1080 resolution.

3 Features

The following is a technical discussion of the key features of *Concrete Earth*, with a focus on its advanced procedural generation. Mathematical models and code snippets are kept to a minimum, edited for clarity rather than accuracy to the original application.

3.1 Procedural Terrain

Given the concept, this game’s terrain is of course intended to evoke a sense of “shunned land” (Trauth et al. 1993). In generating the jutting thorns and blocks so essential to the post-nuclear setting, though, there comes a problem - concavity. Height maps are perfectly adequate for creating rolling hills and valleys, but the fact that each of their (x, z) -coordinates can correspond to only one y -value would prevent any overhangs in *Concrete Earth*’s stark landscape (see Figure 1).




Figure 1: Two sketches of ‘thorny’ terrain; the former can be generated by a height map, whereas the latter is too concave.




Figure 2: Linear partitions of a marching square. With 4 gridpoints, each existing in one of 2 states ('+' for $f_{i,j} \geq c$, '-' for $f_{i,j} < c$), exactly $2^4 = 16$ partitions are possible.




Figure 3: Bounding curve $f(x, y) = 0.16$, at increasing levels of granularity.

Take the two-dimensional analogue of this problem: how does one construct the bounding curve of a (potentially concave, or even disconnected) area? Consider an $n \times n$ square grid, with $n \in \mathbb{N}$,² and let a scalar field $f : [0, 1]^2 \rightarrow \mathbb{R}$ assign each gridpoint (i, j) a corresponding value $f_{i,j} = f(i/n, j/n)$. Figure 2 shows how to partition a cell so as to enclose the gridpoints with $f_{i,j} < c$ constant. Drawing such a partition into every cell of the grid will therefore approximate the *isoline*³ $f(x, y) = c$, with increasing accuracy as $n \rightarrow \infty$.

This is the *marching squares* algorithm, so named for the **for** loop used to ‘march through’ and fill in each of the $n \times n$ square cells. Figure 3, for example, illustrates how it might apply to

$$f(x, y) = \min \left((x - 0.5)^2 + (y - 0.5)^2, 4 \left((x - 0.75)^2 + (y - 0.75)^2 \right) \right),$$

a scalar field describing two overlapping circles.

Given adjacent gridpoints \mathbf{a} , \mathbf{b} such that $f_{\mathbf{a}} < c < f_{\mathbf{b}}$, Figure 2 would have their partitioning line bisect edge AB exactly. However, knowing the partition represents an isoline, this accidentally implies $f(0.5\mathbf{a} + 0.5\mathbf{b}) \approx c$, which may not be true of the scalar field. Using the standard, first-order approximation that f is linear over short distances,

$$f((1-t)\mathbf{a} + t\mathbf{b}) \approx c, \text{ where } t = \frac{c - f_{\mathbf{a}}}{f_{\mathbf{b}} - f_{\mathbf{a}}}.$$

Understanding that it is therefore more accurate for the partition to intersect AB at $(1-t)\mathbf{a} + t\mathbf{b}$, linear interpolation is integrated into the marching squares algorithm (see Figure 4).

¹While the framework registers inputs on press, note that options are only selected *on release*.

²This could equally be an $n \times m$ grid; keeping the dimensions the same is just neater.

³Informally, a ‘contour line’ of the 2D field.

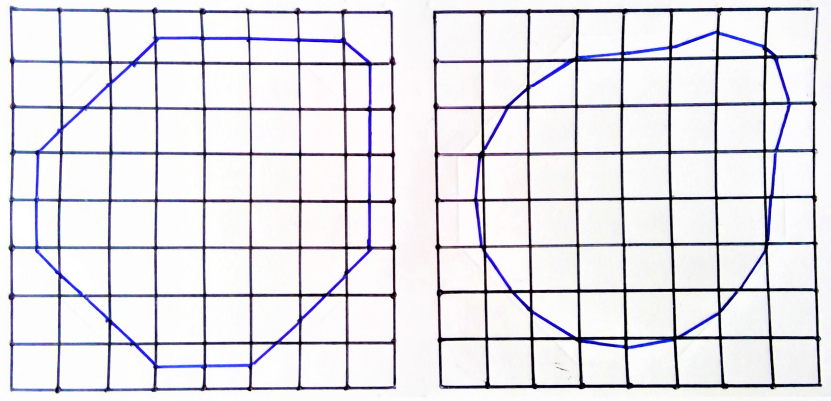


Figure 4: Bounding curve $f(x, y) = 0.16$, before and after linear interpolation.

Returning to the problem at hand, *Concrete Earth* uses *marching cubes* to generate the *isosurfaces* bounding 3D volumes. This requires an $n \times n \times n$ grid, and a new field $f : [0, 1]^3 \rightarrow \mathbb{R}$, the algorithm now filling cubic cells with planar surfaces to partition out the gridpoints (i, j, k) satisfying $f_{i,j,k} < c$ (again using linear interpolation to best approximate the ‘true’ isosurface). Since 3D fields can be hard to visualise, marching cubes are best understood in relation to marching squares. The intuition established above is readily generalised to higher dimensions; whatever the superficial differences, the guiding principle remains the same.

This section will not cover the finer points of the DirectX marching cubes implementation - not how it deconstructs partitioning surfaces into triangles for the vertex/index buffers, nor how it identifies adjacent marching cubes so as to calculate ‘smooth’ vertex normals⁴. The report is only interested in this technique as it relates to *Concrete Earth*’s terrain. Paul Bourke’s *Polygonising a Scalar Field* (1994) remains the definitive source on the algorithm itself.




Figure 5: Planar partitions of a marching cube, before linear interpolation. There are 14 fundamental cases, ignoring symmetries; the code itself considers all $2^8 = 256$ possible orientations.

⁴A weighted average is taken over the surface normals of the surrounding triangular faces. If a face has vertices \mathbf{a} , \mathbf{b} , \mathbf{c} , it will be weighted by $1/|(\mathbf{b} - \mathbf{a}) \times (\mathbf{c} - \mathbf{a})|$, proportional to the inverse of its area.



Figure 6: Isolines of $\mathbf{Hex}(x, y_0, z)$, a horizontal slice of the hexagonal field.

3.1.1 Case Study: Hexes

Having now established a `MarchingCubes` class, generating a landscape becomes a matter of writing the right scalar field. The first step is easy enough: for any isovalue $c \in (0, 1)$,

$$\mathbf{Height}(\mathbf{x}) = y + 0.1 \cdot \mathbf{FBM}(\mathbf{x})$$

models a sweeping expanse, its *fractal brownian motion* component scaled to add roughness whilst avoiding ‘floating islands’. The underlying simplex noise is treated as a black box.

Taking a stylised approach, *Concrete Earth* presents its terrain not as a continuous mass, but a board of isometric, hexagonal tiles. To that extent, it relies equally on the second field

$$\mathbf{Hex}(\mathbf{x}) = \frac{\sqrt{(2x-1)^2 + (2z-1)^2} \cos((\theta \bmod \pi/3) - \pi/6)}{\cos(\pi/6)}, \text{ where } \theta = \arctan2(2z-1, 2x-1).$$

For all the trigonometry, this just describes a vertical hexagonal prism centred at $x = 0.5, z = 0.5$, much the same as $\sqrt{(2x-1)^2 + (2z-1)^2}$ describes a cylinder (see Figure 6).

$$\mathbf{Terrain}_c(\mathbf{x}) = \max(\mathbf{Height}(\mathbf{x}), c \cdot \mathbf{Hex}(\mathbf{x}))$$

then takes a maximum to enclose only those gridpoints that fall below the isovalue in *both* arguments, imposing hexagonal bounds on the existing terrain (the scale factor of c normalising the width of this hex across all isovalues).

3.1.2 Case Study: Landmarks

Physical markers are similarly integrated. To generate one of the afore-mentioned thorns with apex \mathbf{a} , base centred on \mathbf{b} , and half angle ϕ , for instance, the framework defines field

$$\mathbf{Thorn}_{\mathbf{a},\mathbf{b},\phi}(\mathbf{x}) = \frac{1}{\phi} \arccos \left(\frac{(\mathbf{x} - \mathbf{a}) \cdot (\mathbf{b} - \mathbf{a})}{|\mathbf{x} - \mathbf{a}| \cdot |\mathbf{b} - \mathbf{a}|} \right).$$

Now taking a minimum, isosurfaces of

$$\mathbf{Landmarked\ Terrain}_{c;\mathbf{a},\mathbf{b},\phi}(\mathbf{x}) = \min(\mathbf{Terrain}_c(\mathbf{x}), c \cdot \mathbf{Thorn}_{\mathbf{a},\mathbf{b},\phi}(\mathbf{x})),$$

contain all gridpoints below the isovalue in *either* field, adding volume just as a maximum subtracts.⁵

The `Hex` class handling terrain generation introduces a stochastic element, randomising \mathbf{a} , \mathbf{b} , ϕ . Moreover, it might add any number of thorns: taking the minimum over more than two fields just expands the isosurface further. Perhaps the only surprising choice made in procedurally generating these landmarks is keeping their surfaces completely smooth. Rather than simulating weathering and erosion, the conceit of *Concrete Earth* is that these strange, idealised geometries feel distinctly unnatural. The physical markers of shunned land are man-made designs that *should* tap into a sense of the uncanny, in the hopes of communicating to future civilisations that ‘THIS IS NOT A PLACE OF HONOUR’.




Figure 7: *Concrete Earth*’s landmarks. Three designs have been prototyped so far, with more to be added now the terrain generation pipeline is feature-complete.

3.2 Procedural Screen Shaders

The post-processing in *Concrete Earth* is, in many respects, trivial. `vignette_ps.hlsl` demands only two renders-to-texture on every frame: the board itself, and an alpha map of blood vessels sprouting from the edges of the screen. As striking as the final effect is, the shader is surprisingly straightforward in blending the textures into a single, pulsing eye strain overlay; far more worthy of further discussion is how the blood vessels themselves are generated.

⁵These operations do not commute! Applying the maximum first allows thorns that ‘stick out’ from the bounding hex - this is a conscious, if purely aesthetic, choice on the part of the game.

In formal languages, a grammar is a tuple $G = (N, \Sigma, P, \omega_0)$. This contains two disjoint sets of symbols: nonterminals $A, B, \dots \in N$, and terminals $a, b, \dots \in \Sigma$. The production rules in P map nonterminals to strings $\alpha, \beta, \dots \in (N \cup \Sigma)^*$; applied recursively to the axiom $\omega_0 \in (N \cup \Sigma)^*$, these rules can produce increasingly complex *sentences* of terminals and/or nonterminals.⁶

The Chomsky hierarchy (Chomsky 1956) classifies grammars by their production rules:

Type-3. Regular grammars map $A \mapsto a$ or $A \mapsto aB$.

Type-2. Context-free grammars map $A \mapsto \alpha$.

Type-1. Context-sensitive grammars $\alpha A \beta \mapsto \alpha \gamma \beta$.

Type-0. Unrestricted grammars map $\alpha \mapsto \beta$, where α is non-empty.

Note that all Type-3 grammars are also Type-2, all Type-2 grammars also Type-1, and so on.

Suppose, for example, that $N = \{F, G\}$, $\Sigma = \{+, -\}$, $P = \{F \mapsto F + G, G \mapsto F - G\}$, $\omega_0 = F$.

Letting ω_n denote the sentences generated by applying the production rules n times, it follows that

$$\begin{aligned}\omega_1 &= F + G, \\ \omega_2 &= F + G + F - G, \\ \omega_3 &= F + G + F - G + F + G - F - G, \\ \omega_4 &= F + G + F - G + F + G - F - G + F + G + F - G - F + G - F - G, \dots\end{aligned}$$

While these definitions are rather abstract, Lindenmayer (1968) provides a remarkable application. Treating each symbol as an instruction like ‘go forward’ or ‘turn right’, *L-systems* visualise sentences using turtle graphics; when the sentences have been generated recursively by a grammar, the line drawings will inherit their self-similar structure. Consider the above example - reading non-terminals F, G as ‘draw a line while moving 1 unit forwards,’ and terminals \pm as ‘turn by $\pm \pi/2$ on the spot,’ produces fractal *dragon curves* (see Figure 8).




Figure 8: Dragon curves, generated by strings $\omega_2, \omega_4, \dots, \omega_{12}$.

While L-systems are most common in the modelling of 3D plants and other branching structures (Prusinkiewicz & Lindenmayer 1996), *Concrete Earth* only uses them to generate 2D alpha maps. Moreover, it restricts its attention to L-systems paired with context-free grammars.

⁶In mathematical literature, $\omega_0 \in N$ (Hopcroft, Motwani & Ullman 2000), but this project takes an informal approach.

3.2.1 Case Study: Blood Vessels

Generalising the above, *parametric L-systems* act not on symbols A but *modules* $\mathbf{A}(x_1, \dots, x_n)$. By treating non-terminals as structures containing parameters x_1, \dots, x_n , production rules can then be written to act on said parameters.

This project's own `LSystem` class tracks three main values: the length l , width w , and orientation θ with which to draw each line segment. The benefit is, if nothing else, organisational. Even in the basic L-system for dragon curves, terminals \pm are made redundant by the altogether more readable set of production rules

$$\begin{aligned}\mathbf{F}(l, w, \theta) &\mapsto \mathbf{F}(l, w, \theta)\mathbf{G}(l, w, \theta + \pi/2) \\ \mathbf{G}(l, w, \theta) &\mapsto \mathbf{F}(l, w, \theta)\mathbf{G}(l, w, \theta - \pi/2).\end{aligned}$$

Zamir (2001), meanwhile, uses this parametric formulation to visualise bifurcations of blood vessels. Suppose a branch with length l , width w bifurcates into two branches M and m , such that $l_M \geq l_m$. Defining the *asymmetry ratio* $\alpha = l_m/l_M$, it follows that

$$l_M = \frac{l}{(1 + \alpha^3)^{1/3}}, \quad l_m = \frac{\alpha \cdot l}{(1 + \alpha^3)^{1/3}}, \quad w_M = \frac{w}{(1 + \alpha^3)^{1/3}}, \quad w_m = \frac{\alpha \cdot w}{(1 + \alpha^3)^{1/3}}.$$

Furthermore, the branches diverge from their parent at angles

$$\theta_M = \arccos\left(\frac{(1 + \alpha^3)^{4/3} + 1 - \alpha^4}{2(1 + \alpha^3)^{2/3}}\right), \quad \theta_m = \arccos\left(\frac{(1 + \alpha^3)^{4/3} + \alpha^4 - 1}{2\alpha^2(1 + \alpha^3)^{2/3}}\right).$$

This application is therefore capable of reproducing Zamir's results using an L-system with the single production rule $\mathbf{C}(l, w, \theta) \mapsto \mathbf{X}(l, w, \theta)[\mathbf{C}(l_M, w_M, \theta + \theta_M)]\mathbf{C}(l_m, w_m, \theta - \theta_m)$ (see Figure 9).




Figure 9: Zamir's model of arterial branching, generated with asymmetry ratios $\alpha = 1.0, 0.8, \dots, 0.2$. This is an explicit copy of Zamir (2001, Figure 8), recreated within the DirectX framework.

Liu et al. (2010) expand on this research with a further stochastic component - they assign *multiple* production rules to a single non-terminal, choosing which to apply using a weighted probability

distribution. *Concrete Earth* incorporates such randomness into its own rules for blood vessels:

$$\begin{array}{lcl}
\mathbf{C}(l, w, \theta) & \xrightarrow{0.45} & \mathbf{X}(l, w, \theta)[\mathbf{L}(l_M, w_M, \theta + \theta_M)]\mathbf{R}(l_m, w_m, \theta - \theta_m) \\
\mathbf{C}(l, w, \theta) & \xrightarrow{0.45} & \mathbf{X}(l, w, \theta)[\mathbf{L}(l_m, w_m, \theta + \theta_m)]\mathbf{R}(l_M, w_M, \theta - \theta_M) \\
\mathbf{C}(l, w, \theta) & \xrightarrow{0.10} & \mathbf{X}(l, w, \theta)\mathbf{C}(l, w, \theta) \\
\mathbf{L}(l, w, \theta) & \xrightarrow{1.00} & \mathbf{X}(l, w, \theta)\mathbf{C}(l_M, w_M, \theta - \theta_M) \\
\mathbf{R}(l, w, \theta) & \xrightarrow{1.00} & \mathbf{X}(l, w, \theta)\mathbf{C}(l_M, w_M, \theta + \theta_M)
\end{array}$$

These describe a capillary with a 45% chance of bifurcating with branch M tacking clockwise, a 45% chance of bifurcating with M tacking anticlockwise, and a 10% chance of extending forwards without any branching. The deterministic production rules on \mathbf{L} , \mathbf{R} provide course correction, ensuring the models can be pointed towards the centre of the screen without veering off in either direction. Further informal tweaks can be found in the `LBloodVessel` class itself.

Finally, assigning each `line` segment a (strictly positive) `depth` value and taking the `maxDepth` across all such segments, a normalised `intensity` factor scales the model as follows:

```

for (int i = 0; i < line.size(); i++)
{
    float growth = 1.0f-(1.0f-intensity)*maxDepth/line[i].depth;
    line[i].length *= std::max(0.0f, std::min(growth, 1.0f));
}

```

As `intensity` increases from `0.0f` to `1.0f`, the blood vessels grow, the ‘deepest’ branches emerging first.⁷ Rendering these to texture results in a visceral, procedurally-animated post-process effect, more responsive than passing in a static alpha map.




Figure 10: *Concrete Earth*’s stochastic model of arterial branching, generated with different seeds.

3.3 Procedural Narrative

Concrete Earth was originally conceived as a work of interactive fiction, a mix of authorship and procedural text generation. Only a fraction of the content has been written so far, so this section will focus on the *intended* narrative, and the infrastructure already in place to facilitate it.

⁷On a more technical level, *Concrete Earth* treats a branch’s `depth` as a function of the `lengths` of its child branches. Though it came about from a typo in the code, the definition stuck - by delaying growth around bifurcation points, it achieves a nice ‘sprouting’ effect!

3.3.1 Grammars

Given their origin in linguistics, it is perhaps unsurprising that grammars (see Section 3.2) find use in the field of procedural narrative - *Tracery* (Compton et al. 2015) being a prime example. This is, formally, a stochastic, context-free grammar, with uniformly-distributed production rules for each non-terminal; it iterates a given axiom until all non-terminals (demarcated by {BRACE} delimiters) have been replaced. To bemoan the generator as better suited to Twitter bots than interactive storytelling is to misunderstand its design. Compton et al. present a lightweight tool that is left *deliberately* narrow, intended for ease-of-use by even the most fledgling authors.

More substantial works, then, adapt *Tracery*’s grammar-based approach to their own specifications. In *The Annals of the Parrigues*, Short uses a consistent, context-sensitive generator that “defines any facts about the world that aren’t already defined at the moment of generation” (2015, p. 83). *Voyageur* (Dias 2018) attaches weightings to the distributions of its production rules (Dias 2019). Producing storytelling is an art form; as much as *Concrete Earth* wears the above influences on its sleeve, to try and consolidate these into an ‘optimal’, one-size-fits-all generator would be missing the point. This project therefore uses an implementation of *Tracery* with its own custom modifications.

Recency To reduce repetition of a word or phrase, one might consider when it has last been used. The custom grammar treats this as a rank: given a non-terminal with N possible production rules, the most recently applied rule having a *recency*⁸ of $N - 1$, the second most recent $N - 2$, and so on down to the rule that was applied longest ago (or indeed, never!), with recency 0. Though it does not factor in the intervals between uses, nor the second, third, etc, most recent use, there is no need to overcomplicate a metric that already achieves the desired effect (Kazemi 2019).

These ranks are straightforwardly integrated into the production rules’ probability distribution. The likelihood of a rule with recency n is scaled by $p^{n/(N-1)}$, such that (all other weightings being equal) the most recently used rule will be p times as likely as the least recent.⁹

Consistency Intended to create a feeling of transience, landmarks in *Concrete Earth* can be visited once and once only, and as such do not need logged in an overall ‘world model’. The game’s approach to characterisation is not so simple. With a narrative largely centred around picking up passengers, journeying with them, then parting ways, it should follow that if a detail about a party member has already been established, the text generator can consistently recover it.

The grammar’s `GenerateSentence` function will therefore accept (up to) two pointers to `StoryCharacter` structs. These are handled by the *consistency markers* `*`, `**`, `?`:

- `{*A}` tells the grammar to check the first character argument for key `A`. If a corresponding value exists, it will be substituted directly in place of `{*A}`; if not, the grammar will choose a production rule as usual, storing this choice in the character struct for future recall.
- `{**A}` acts similarly on the second character argument.
- `{?A}` inherits the consistency marker from its preceding non-terminal, allowing for ‘follow-up’ information. If a `{**SURNAME}` is replaced by `{?BARREL #1}-{?BARREL #2}`, for instance, both halves of the second character’s double-barrelled surname will also need to be consistent.¹⁰

Conditionality By their very nature, grammars struggle “to encode high level relationships between different things generated at different points” in a sentence (Compton 2019). The game

⁸Or ‘dryness’, as *Improv*, the generator for *Voyageur*, would call it (Dias 2019).

⁹The same string `alpha` might replace two different non-terminals `{A}`, `{B}`. Recognising these as distinct production rules, the recency with which one was applied to `{A}` has no bearing on when the other is applied to `{B}`, or vice versa.

¹⁰Note that `?` can also inherit an empty string, if the preceding non-terminal is not being remembered.

tries to work around this with *nested non-terminals*: in receiving a term $\{\{A\} B\}$, the grammar will first replace $\{A\}$ with some **alpha**, creating a new non-terminal $\{\text{alpha } B\}$ dependent on the initial choice. While this solution doesn't exactly scale,¹¹ it is an author-friendly means of tracking key conditions.

Consolidating the above, consider how *Concrete Earth* might use the Russian naming convention

$\{*\{*\text{GENDER}\} \text{FORENAME}\} \{*\{*\text{GENDER}\} \text{PATRONYMIC}\} \{*\text{SURNAME}\}$

to introduce a new character:

1. $\{*\{*\text{GENDER}\} \text{FORENAME}\} \mapsto \{*\text{MASCULINE} \text{FORENAME}\}$: The grammar starts with the nested $\{*\text{GENDER}\}$ non-terminal. The asterisk tells it to use the character's established gender, but - finding nothing - it randomly decides whether they are a man, woman, or non-binary, passing this value into sentence and the `StoryCharacter` struct alike.
2. $\{*\text{MASCULINE} \text{FORENAME}\} \mapsto \text{LEO}$: A forename is chosen according to gender identity; even if that gender identity is never made explicit to the player, internal consistency is essential.
3. $\{*\{*\text{GENDER}\} \text{PATRONYMIC}\} \mapsto \{*\text{MASCULINE} \text{PATRONYMIC}\}$: Now that a value for `GENDER` has been established, it is substituted back in to the second non-terminal.
4. $\{*\text{MASCULINE} \text{PATRONYMIC}\} \mapsto \text{NIKOLAYEVICH}$: The character's forename is paired with an appropriate patronymic ("Nikolayevich" here meaning "Son of Nikolay").¹²
5. $\{*\text{SURNAME}\} \mapsto \text{TOLSTOY}$: The final non-terminal comes with no conditions, and is replaced outright. This, too, is saved within the character struct.

In this case, the grammar terminates after a single iteration. Should $\{*\text{BARREL } \#1\} - \{*\text{BARREL } \#2\}$ have replaced $\{*\text{SURNAME}\}$, however, the new non-terminals would trigger a second parse.

The framework might feel too limited, from a dramatic perspective. Personality is not immutable; for all this focus on consistency, good storytelling is just as reliant on character *development*. Indeed, in a piece of interactive fiction one might expect to have an actual impact on the characters' lives. The grammar surely requires some sort of override...

3.3.2 Storylets

Storylets (Kreminski & Wardrip-Fruin 2018) are narrative modules, best thought of as the largest 'unit of story' in a given project. The definition is broad (again, this is not an exact science), but storylets are generally understood as a dynamic, yet controlled, option for storytelling. They come with their own internal structures, which may well be parsed by grammars or other text generators. What's more, being highly reorderable, they are best surfaced by a surrounding *content selection architecture* - it is common to use salience scores or a weighted distribution to prioritise the more surprising modules when their conditions are met.

In a full release of *Concrete Earth*, the player will first be presented with a scene (the 'beginning' of the storylet) and an accompanying set of choices (the 'middle') that fit with the world's current state. When they select one of the options, the architecture will find a fitting consequence (the 'end'). It is a streamlined internal structure, with the beginning, middle and end *all* capable of updating the world model, making changes beyond just filling in uninitialised character details.

¹¹Handling all possible cases of the nested non-terminal $\{\{A1\} \{A2\} \{A3\} \{A4\} \{A5\} B\}$ will get out of hand!

¹²As far as this report can tell, there are no gender-neutral patronymics in Russian - without actually speaking the language, it seems most appropriate to drop these from non-binary characters' names entirely. Cook (2019) offers an invaluable discussion of representation within procedural narrative, and how generators can (whether by accident or authorial intent) reveal prejudices within the societies they create.

The current implementation, however, is limited. The content selection architecture is trivial, surfacing storylets at random (so long as the world model satisfies their minimal sets of conditions). Moreover, the player's effect on the narrative amounts to an arbitrary choice of whether or not to add a given passenger to their party. In the future, the storylet framework could be essential in tracing out meaningful character arcs; for now, it exists only in support of other, more successful features.

4 Code Organisation

This application is, amongst other things, a continuation of previous work. Being built on my framework for CMP502, *Concrete Earth* has been a chance to re-evaluate much of the prior code. Shaders, for instance, were once handled with an unwieldy number of C++ classes: a `LightShader` that passed lighting details into the pixel shader, a `SkyboxShader` passing in an environment map, etc. These have now been folded back into their parent `Shader` class, with `public` functions to turn on the relevant buffers. Though it might seem very basic, identifying and acting on such bad practices early on has been essential to the overall development process.

As identified above, procedural modelling comes with two key challenges - establishing general, multi-purpose frameworks, then tailoring these to specific use cases. The code has been organised to reflect this dichotomy. Landmark `Fields` are constructed independently of the `MarchingCubes`; while an `LSystems` class uses formal grammars to generate turtle drawings, it is through child classes like `LBloodVessel` and `LDragonCurve` that the grammars are actually described. Once initialised, `Game.cpp` treats these procedural models as it would any others, using a choice of vertex and pixel shader to render them within the graphics pipeline.

More than other DirectX assignments, this project has come with a need for genuine gameplay. A `Board` class has therefore been included to handle the central loop. At a very conceptual level, it exists as a mediator, reconciling the player's physical position on a landscape of `Hexes` with the more ethereal, text-based mechanics.

A narrative pipeline emerges. On every move, the `Board` checks if it should start a new scene. If so, it passes the player's current location (whether they are near thorns, a towering monolith, etc) down to a `StoryEngine`, which then surfaces a new storylet through its `Architecture` variable. Its beginning and middle sections are parsed by a `Grammar`, and passed back upwards. When the `Board` then tells the `StoryEngine` the player's choice, it similarly receives an ending. Note that the `Architecture` and `Grammar` will both need access to a `StoryWorld` struct, the `StoryEngine` continuously updating one of these as its definitive world model.

Storylets are written to screen using the `ImGui` library - the UI is minimal, handled almost entirely within a single `Game::SetupGUI` function. This application also relies on `nlohmann's JSON API`.¹³ The `font` and `Geiger counter SFX` were sourced online; all code is cited as appropriate.

Finally, *Concrete Earth* was written with good coding practices in mind. Functions and variables follow a consistent naming convention, focused on readability. The DRY principle has been followed, at least within reason. Even the current, informal use of comments proves effective, these breaking down the logic behind key snippets, highlighting potential bugs, and overall acting as incisive memos-to-self¹⁴ (a collaborative project would obviously require a higher standard of documentation). GitHub was used for version control.

¹³Indeed, it is a conscious organisational choice to store word banks of terminals and non-terminals within their own self-contained JSON files, rather than hardcode them into the `Grammar` itself.

¹⁴This is notably untrue of `Game.cpp`. All figures in this report have been expressly produced for CMP505 - though unused in the final application, the code used to render these figures remains, left commented out in large, unclear blocks for the sake of posterity.

5 Evaluation

5.1 Features

If one aspect of this project constitutes an unambiguous success, it is surely the use of L-systems. *Concrete Earth* provides not only a robust, parametric framework, but also a genuinely interesting application thereof. While the game’s models of blood vessels are perfectly acceptable in themselves, it is their integration with the rest of the post-process pipeline that (literally) allows them to shine. Not that there isn’t more to do, of course - the procedural animation, for example, still feels rather rudimentary - but the `LSystem` class is in a good place. Above all else, it is *versatile*, a foundational piece of code to be transposed into any number of future projects.

The role of marching cubes, by contrast, feels much more narrow: height maps have a well-known limitation, to which this algorithm is the well-known solution (see Section 3.1). Implementing the code was a significant technical challenge, but largely successful, and though the generation process does not run especially quickly¹⁵ a potential solution has already been identified. Bourke’s discussion of marching cubes is accompanied by an introduction to *marching tetrahedrons* (1997), a variation of the algorithm that would hopefully improve the efficiency of the code.¹⁶ Working with a tetrahedral grid will come with additional nuances (e.g. cell orientation), but I now feel I’ve built enough intuition to take on this more advanced technique.

The more interesting question here is how effectively the marching cubes enhance *Concrete Earth*’s setting. As much as they are used to procedurally generate terrain, that terrain can feel somewhat lacking - while this report has already acknowledged the relatively few landmarks available, it’s no less notable that they are left untextured. With more time, I would have been keen to write a pixel shader to handle this, but having already experimented with procedural texturing in CMP502 it felt more fruitful to focus on new techniques. Note also that all terrain hexes are generated at runtime; updating these dynamically would be a relatively minor change that greatly increases variety.



Figure 11: 3D Voronoi noise. Experiments with voxel textures did not make it past initial prototyping.

To the extent that the game has any single ‘special feature’, it would surely be its approach to procedural storytelling. On a technical level, the `Grammar` class works exactly as intended: it offers all the functionality of *Tracery*, with targeted tweaks. For all my interest in narrative frameworks,

¹⁵Hence why the executable will often need a few minutes to ‘boot up’.

¹⁶After all, where there are 256 possible partitions of a cube, tetrahedrons only come with $2^4 = 16$...

though, I am fundamentally not a writer, a fact reflected by a rather limited collection of corpora (plus a typo here and there!). Just like the marching cubes, the grammar lacks the *content* to reach its full potential.

5.2 Code Organisation

This feeds into a deeper, organisational issue. Having known how few storylets would be submitted within this prototype, the accompanying infrastructure ended up overly niche, with large sections of `StoryEngine.cpp` and `Architecture.cpp` hardcoded around the contents of `Storylets.json`. The code *works*, certainly, but in terms of how much I'll need to refactor it is functionally useless. For now, it would be most valuable to take a step back, and draft a more representative sample of storylets: more content will inevitably inform a more general approach to the programming.

The rest of the application is far more robust - indeed, the above seems like a very singular point of failure when taken in its broader context. Coming into the project with previous experience in DirectX, I was determined to use create a much more streamlined experience this time around. It's been a valuable philosophy, the benefits extending far beyond my own, personal workflow. Given the ambitious nature of *Concrete Earth*, that the application performs well enough on even a low-end laptop should speak to the effectiveness of the organisational strategy.

6 Conclusions

, has of course . . . Approaching text generation as an art rather than a science, for instance, was essential to understanding my own narrative pipeline. I found L-systems easier to organise once I understood the underlying theory, and could clearly delineate between the grammar and turtle drawing components. Indeed, I'd have got to grips with marching cubes far more quickly had I started out in two dimensions! These are all valuable observations, if somewhat surface-level.

A more meaningful throughline of this report might be the relationship between implementing features, and generating content *with* those features. As much as I , I'm not so sure I'd draw that same distinction. ; conversely, The eL-system; conversely, my content selection architecture feels so

[Last: iteration on CMP502...]

References

- Bourke, P. (1994), ‘Polygonising a Scalar Field’, Available at: <http://paulbourke.net/geometry/polygonise/>. (Accessed: 9 February 2023).
- Bourke, P. (1997), ‘Polygonising a Scalar Field Using Tetrahedrons’, Available at: <http://paulbourke.net/geometry/polygonise/>. (Accessed: 9 February 2023).
- Chomsky, N. (1956), ‘Three Models for the Description of Language’, *IRE Transactions on Information Theory* **2**(3), 113–124.
- Compton, K. (2019), ‘Getting Started with Generators’, in T. Short & T. Adams, eds, ‘*Procedural Storytelling in Game Design*’, 2nd edn, Boca Raton, FL, USA: Taylor & Francis, pp. 3–17.
- Compton, K., Kybartas, B. & Mateas, M. (2015), Tracery: An Author-Focused Generative Text Tool, in ‘*8th International Conference on Interactive Digital Storytelling*’, Copenhagen, Denmark: 30 November–4 December, pp. 154–161.
- Cook, M. (2019), ‘Ethical Procedural Generation’, in T. Short & T. Adams, eds, ‘*Procedural Storytelling in Game Design*’, 2nd edn, Boca Raton, FL, USA: Taylor & Francis, pp. 49–62.
- Dias, B. (2018), *Voyageur*, self-published.
- Dias, B. (2019), ‘Procedural Descriptions in *Voyageur*’, in T. Short & T. Adams, eds, ‘*Procedural Storytelling in Game Design*’, 2nd edn, Boca Raton, FL, USA: Taylor & Francis, pp. 193–207.
- Hopcroft, J., Motwani, R. & Ullman, J. D. (2000), *Introduction to Automata Theory, Languages, and Computation*, 2nd edn, Boston, MA, USA: Addison-Wesley.
- Kazemi, D. (2019), ‘Keeping Procedural Generation Simple’, in T. Short & T. Adams, eds, ‘*Procedural Storytelling in Game Design*’, 2nd edn, Boca Raton, FL, USA: Taylor & Francis, pp. 17–22.
- Kreminski, M. & Wardrip-Fruin, N. (2018), Sketching a Map of the Storylets Design Space, in ‘*11th International Conference on Interactive Digital Storytelling*’, Dublin, Ireland: 5–8 December, pp. 160–164.
- Lindenmayer, A. (1968), ‘Mathematical Models for Cellular Interactions in Development II. Simple and Branching Filaments With Two-Sided Inputs’, *Journal of Theoretical Biology* **18**(3), 300–315.
- Liu, X., Liu, H., Hao, A. & Zhao, Q. (2010), Simulation of Blood Vessels for Surgery Simulators, in ‘*2010 International Conference on Machine Vision and Human-machine Interface*’, pp. 377–380.
- Prusinkiewicz, P. & Lindenmayer, A. (1996), *The Algorithmic Beauty of Plants*, 2nd edn, Berlin, Germany: Springer-Verlag.
- Short, E. (2015), *The Annals of the Parrigues*, self-published.
- Trauth, K., Hora, S. & Guzowski, R. (1993), Expert Judgment on Markers to Deter Inadvertent Human Intrusion Into the Waste Isolation Pilot Plant, Technical report, SANDIA National Laboratories, Albuquerque.
- Zamir, M. (2001), ‘Arterial Branching Within the Confines of Fractal L-System Formalism’, *The Journal of General Physiology* **118**, 267–276.